



# **ETC3550**

## **Applied forecasting for business and economics**

Ch3. Evaluating modelling accuracy

[OTexts.org/fpp3/](https://OTexts.org/fpp3/)

# Outline

- 1 Residual diagnostics
- 2 Evaluating forecast accuracy
- 3 Time series cross-validation

# Outline

- 1 Residual diagnostics
- 2 Evaluating forecast accuracy
- 3 Time series cross-validation

# Fitted values

- $\hat{y}_{t|t-1}$  is the forecast of  $y_t$  based on observations  $y_1, \dots, y_t$ .
- We call these “fitted values”.
- Sometimes drop the subscript:  $\hat{y}_t \equiv \hat{y}_{t|t-1}$ .
- Often not true forecasts since parameters are estimated on all data.

## For example:

- $\hat{y}_t = \bar{y}$  for average method.
- $\hat{y}_t = y_{t-1} + (y_T - y_1)/(T - 1)$  for drift method.

# Forecasting residuals

**Residuals in forecasting:** difference between observed value and its fitted value:  $e_t = y_t - \hat{y}_{t|t-1}$ .

# Forecasting residuals

**Residuals in forecasting:** difference between observed value and its fitted value:  $e_t = y_t - \hat{y}_{t|t-1}$ .

## Assumptions

- 1  $\{e_t\}$  uncorrelated. If they aren't, then information left in residuals that should be used in computing forecasts.
- 2  $\{e_t\}$  have mean zero. If they don't, then forecasts are biased.

# Forecasting residuals

**Residuals in forecasting:** difference between observed value and its fitted value:  $e_t = y_t - \hat{y}_{t|t-1}$ .

## Assumptions

- 1  $\{e_t\}$  uncorrelated. If they aren't, then information left in residuals that should be used in computing forecasts.
- 2  $\{e_t\}$  have mean zero. If they don't, then forecasts are biased.

## Useful properties (for prediction intervals)

- 3  $\{e_t\}$  have constant variance.
- 4  $\{e_t\}$  are normally distributed.

# Example: Google stock price


```
google_2015 <- tsibbledata::gafa_stock %>%  
  filter(Symbol == "GOOG", year(Date) == 2015) %>%  
  mutate(trading_day = row_number()) %>%  
  update_tsibble(index = trading_day, regular = TRUE)
```

```
## # A tsibble: 252 x 9 [1]  
## # Key:      Symbol [1]  
##   Symbol Date      Open  High   Low Close  
##   <fct>  <date>    <dbl> <dbl> <dbl> <dbl>  
## 1 GOOG   2015-01-02   526.  528.  521.  522.  
## 2 GOOG   2015-01-05   520.  521.  510.  511.  
## 3 GOOG   2015-01-06   512.  513.  498.  499.  
## 4 GOOG   2015-01-07   504.  504.  497.  498.  
## 5 GOOG   2015-01-08   495.  501.  488.  500.  
## 6 GOOG   2015-01-09   502.  502.  492.  493.  
## 7 GOOG   2015-01-12   492.  493.  485.  490.  
## 8 GOOG   2015-01-13   496.  500.  490.  493.  
## 9 GOOG   2015-01-14   492.  500.  490.  498.  
## 10 GOOG  2015-01-15   503.  503.  495.  499.
```



# Example: Google stock price

```
google_2015 %>%  
  autoplot(Close) +  
  xlab("Day") + ylab("Closing Price (US$)") +  
  ggtitle("Google Stock (daily ending 6 December 2013)")
```



# Example: Google stock price

Naïve forecast:

$$\hat{y}_{t|t-1} = y_{t-1}$$

# Example: Google stock price

Naïve forecast:

$$\hat{y}_{t|t-1} = y_{t-1}$$

$$e_t = y_t - y_{t-1}$$

# Example: Google stock price

Naïve forecast:


$$\hat{y}_{t|t-1} = y_{t-1}$$

$$e_t = y_t - y_{t-1}$$

Note:  $e_t$  are one-step-forecast residuals


# Example: Google stock price

```
fit <- google_2015 %>% model(NAIVE(Close))  
augment(fit) %>%  
  ggplot(aes(x = trading_day)) +  
  geom_line(aes(y = Close, colour = "Data")) +  
  geom_line(aes(y = .fitted, colour = "Fitted")) +  
  xlab("Day") + ylab("Closing Price (US$)") +  
  ggtitle("Google Stock (daily ending 6 December 2013)")
```




# Example: Google stock price

```
augment(fit) %>%  
  autoplot(.resid) + xlab("Day") + ylab("") +  
  ggtitle("Residuals from naïve method")
```




# Example: Google stock price

```
augment(fit) %>%  
  ggplot(aes(x = .resid)) +  
  geom_histogram(bins = 30) +  
  ggtitle("Histogram of residuals")
```



# Example: Google stock price

```
residuals(fit) %>% ACF(.resid) %>%  
  autoplot() + ggtitle("ACF of residuals")
```





# ACF of residuals

- We assume that the residuals are white noise (uncorrelated, mean zero, constant variance). If they aren't, then there is information left in the residuals that should be used in computing forecasts.
- So a standard residual diagnostic is to check the ACF of the residuals of a forecasting method.
- We *expect* these to look like white noise.

## Portmanteau tests

Consider a *whole set* of  $r_k$  values, and develop a test to see whether the set is significantly different from a zero set.

# Portmanteau tests

Consider a *whole* set of  $r_k$  values, and develop a test to see whether the set is significantly different from a zero set.

## Box-Pierce test

$$Q = T \sum_{k=1}^h r_k^2$$

where  $h$  is max lag being considered and  $T$  is number of observations.

- If each  $r_k$  close to zero,  $Q$  will be **small**.
- If some  $r_k$  values large (positive or negative),  $Q$  will be **large**.

# Portmanteau tests

Consider a *whole* set of  $r_k$  values, and develop a test to see whether the set is significantly different from a zero set.

## Ljung-Box test

$$Q^* = T(T+2) \sum_{k=1}^h (T-k)^{-1} r_k^2$$

where  $h$  is max lag being considered and  $T$  is number of observations.

- My preferences:  $h = 10$  for non-seasonal data,  $h = 2m$  for seasonal data.
- Better performance, especially in small samples.

# Portmanteau tests

- If data are WN,  $Q^*$  has  $\chi^2$  distribution with  $(h - K)$  degrees of freedom where  $K$  = no. parameters in model.
- When applied to raw data, set  $K = 0$ .
- For the Google example:

```
# lag=h and fitdf=K
```

```
Box.test(augment(fit)$resid, lag=10, fitdf=0, type="Lj")
```

```
##
```

```
## Box-Ljung test
```

```
##
```

```
## data: augment(fit)$resid
```

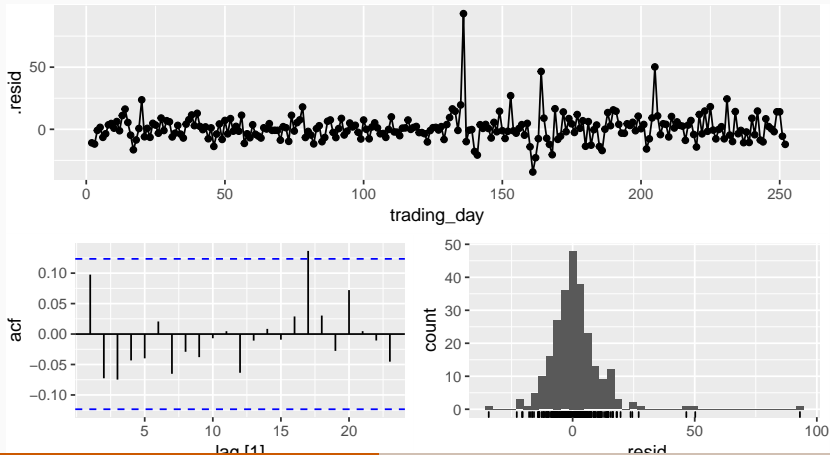
```
## X-squared = 7.9141, df = 10, p-value =
```

```
## 0.6372
```

# gg\_tsdisplay function

```
augment(fit) %>%
```

```
gg_tsdisplay(.resid, plot_type = "histogram")
```



## Your turn

Compute seasonal naïve forecasts for quarterly Australian beer production from 1992.

```
recent <- aus_production %>% filter(year(Time)  
fit <- recent %>% model(snaive(Beer))  
fit %>% forecast() %>% autoplot(recent)
```

Test if the residuals are white noise.

```
Box.test(augment(fit)$resid, lag=10, fitdf=0,  
augment(fit) %>% gg_tsdisplay(.resid, plot_type
```

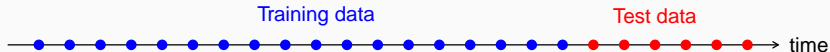
What do you conclude?

# Outline

- 1 Residual diagnostics
- 2 Evaluating forecast accuracy
- 3 Time series cross-validation



# Training and test sets



- A model which fits the training data well will not necessarily forecast well.
- A perfect fit can always be obtained by using a model with enough parameters.
- Over-fitting a model to data is just as bad as failing to identify a systematic pattern in the data.
- The test set must not be used for *any* aspect of model development or calculation of forecasts.
- Forecast accuracy is based only on the test set.

# Forecast errors


Forecast “error”: the difference between an observed value and its forecast.

$$e_{T+h} = y_{T+h} - \hat{y}_{T+h|T},$$

where the training data is given by  $\{y_1, \dots, y_T\}$

- Unlike residuals, forecast errors on the test set involve multi-step forecasts.
- These are *true* forecast errors as the test data is not used in computing  $\hat{y}_{T+h|T}$ .

# Measures of forecast accuracy



# Measures of forecast accuracy

$y_{T+h}$  =  $(T + h)$ th observation,  $h = 1, \dots, H$

$\hat{y}_{T+h|T}$  = its forecast based on data up to time  $T$ .

$$e_{T+h} = y_{T+h} - \hat{y}_{T+h|T}$$

$$\text{MAE} = \text{mean}(|e_{T+h}|)$$

$$\text{MSE} = \text{mean}(e_{T+h}^2)$$

$$\text{RMSE} = \sqrt{\text{mean}(e_{T+h}^2)}$$

$$\text{MAPE} = 100\text{mean}(|e_{T+h}|/|y_{T+h}|)$$

# Measures of forecast accuracy

$y_{T+h}$  =  $(T + h)$ th observation,  $h = 1, \dots, H$

$\hat{y}_{T+h|T}$  = its forecast based on data up to time  $T$ .

$$e_{T+h} = y_{T+h} - \hat{y}_{T+h|T}$$

$$\text{MAE} = \text{mean}(|e_{T+h}|)$$

$$\text{MSE} = \text{mean}(e_{T+h}^2)$$

$$\text{RMSE} = \sqrt{\text{mean}(e_{T+h}^2)}$$

$$\text{MAPE} = 100\text{mean}(|e_{T+h}|/|y_{T+h}|)$$

- MAE, MSE, RMSE are all scale dependent.
- MAPE is scale independent but is only sensible if  $y_t \gg 0$  for all  $t$ , and  $y$  has a natural zero.

# Measures of forecast accuracy

## Mean Absolute Scaled Error

$$\text{MASE} = \text{mean}(|e_{T+h}|/Q)$$

where  $Q$  is a stable measure of the scale of the time series  $\{y_t\}$ .

Proposed by Hyndman and Koehler (IJF, 2006).

For non-seasonal time series,

$$Q = (T - 1)^{-1} \sum_{t=2}^T |y_t - y_{t-1}|$$

works well. Then MASE is equivalent to MAE relative to a naïve method.

# Measures of forecast accuracy

## Mean Absolute Scaled Error

$$\text{MASE} = \text{mean}(|e_{T+h}|/Q)$$

where  $Q$  is a stable measure of the scale of the time series  $\{y_t\}$ .


Proposed by Hyndman and Koehler (IJF, 2006).

For seasonal time series,

$$Q = (T - m)^{-1} \sum_{t=m+1}^T |y_t - y_{t-m}|$$

works well. Then MASE is equivalent to MAE relative to a seasonal naïve method.

# Measures of forecast accuracy





# Training set accuracy

```
recent_production <- aus_production %>%  
  filter(year(Quarter) >= 1992)  
train <- recent_production %>% filter(year(Quarter) <= 2007)  
beer_fit <- train %>%  
  model(  
    Mean = MEAN(Beer),  
    Naïve = NAIVE(Beer),  
    Seasonal naïve = SNAIVE(Beer),  
    Drift = RW(Beer ~ drift())  
  )
```

	RMSE	MAE	MAPE	MASE
Mean method	43.62858	35.23438	7.886776	2.463942
Naïve method	65.31511	54.73016	12.164154	3.827284
Seasonal naïve method	16.78193	14.30000	3.313685	1.000000
Drift method	65.31337	54.76795	12.178793	3.829927

# Test set accuracy

```
beer_fc <- beer_fit %>%  
  forecast(h = 10)  
accuracy(beer_fc, recent_production)
```

	RMSE	MAE	MAPE	MASE
Drift method	64.90129	58.87619	14.577487	4.1172161
Mean method	38.44724	34.82500	8.283390	2.4353147
Naïve method	62.69290	57.40000	14.184424	4.0139860
Seasonal naïve method	14.31084	13.40000	3.168503	0.9370629

## Poll: true or false?

- 1 Good forecast methods should have normally distributed residuals.
- 2 A model with small residuals will give good forecasts.
- 3 The best measure of forecast accuracy is MAPE.
- 4 If your model doesn't forecast well, you should make it more complicated.
- 5 Always choose the model with the best forecast accuracy as measured on the test set.

# Outline

- 1 Residual diagnostics
- 2 Evaluating forecast accuracy
- 3 Time series cross-validation


# Time series cross-validation

## Traditional evaluation




# Time series cross-validation

## Traditional evaluation



## Time series cross-validation




# Time series cross-validation

## Traditional evaluation



## Time series cross-validation



- Forecast accuracy averaged over test sets.
- Also known as “evaluation on a rolling forecasting origin”

# Creating the rolling training sets

There are three main rolling types which can be used.

- Stretch: extends a growing length window with new data.
- Slide: shifts a fixed length window through the data.
- Tile: moves a fixed length window without overlap.

Three functions to roll a tsibble: `stretch_tsibble()`, `slide_tsibble()`, and `tile_tsibble()`.

For time series cross-validation, stretching windows are most commonly used.



# Time series cross-validation

Stretch with a minimum length of 3, growing by 1 each step.

```
google_2015_stretch <- google_2015 %>%  
  stretch_tsibble(.init = 3, .step = 1) %>%  
  filter(.id != max(.id))
```

```
## # A tsibble: 31,623 x 4 [1]  
## # Key:      .id [249]  
##   Date      Close trading_day  .id  
##   <date>    <dbl>          <int> <int>  
## 1 2015-01-02  522.              1      1  
## 2 2015-01-05  511.              2      1  
## 3 2015-01-06  499.              3      1  
## 4 2015-01-02  522.              1      2  
## # ... with 3.162e+04 more rows
```

# Time series cross-validation

Estimate RW w/ drift models for each window.

```
fit_cv <- google_2015_stretch %>%  
  model(RW(Close ~ drift()))
```

```
## # A mable: 249 x 3  
## # Key:      .id, Symbol [249]  
##      .id Symbol RW(Close ~ drift())  
##    <int> <fct>  <model>  
## 1      1 GOOG   <RW w/ drift>  
## 2      2 GOOG   <RW w/ drift>  
## 3      3 GOOG   <RW w/ drift>  
## 4      4 GOOG   <RW w/ drift>  
## # ... with 245 more rows
```

# Time series cross-validation

Produce one step ahead forecasts from all models.

```
fc_cv <- fit_cv %>%  
  forecast(h=1)
```

```
## # A tibble: 249 x 5  
## #   .id Symbol trading_day Close  
##   <int> <fct>         <int> <dbl>  
## 1     1  G00G             4  488.  
## 2     2  G00G             5  490.  
## 3     3  G00G             6  494.  
## 4     4  G00G             7  488.  
## # ... with 245 more rows, and 1 more  
## #   variable: .distribution <dist>
```

# Time series cross-validation

```
# Cross-validated  
fc_cv %>% accuracy(google_2015)  
# Training set  
google_2015 %>% model(NAIVE(Close)) %>% accuracy()
```

	RMSE	MAE	MAPE
Cross-validation	11.26819	7.261240	1.194024
Training	11.18958	7.127985	1.170985

A good way to choose the best forecasting model is to find the model with the smallest RMSE computed using time series cross-validation.